Thesis Structure

Central Bank Dissonance and Asset Returns

Hierarchical Outline · Draft v0.1

Hierarchical outline with three levels of headings. Word counts are indicative for a master's-scale deliverable; the doctoral upgrade adds depth within the same architecture rather than new chapters.

Total body word count. ≈ 47,000 words for master's. ≈ 70,000–80,000 words for doctoral upgrade.

Chapter count. Nine substantive chapters plus front matter and back matter (including a timestamped pre-registration appendix).

Front Matter

- Title page
- Declaration of originality
- Abstract (300 words)
- Acknowledgements
- Table of contents
- List of figures
- List of tables
- List of abbreviations
- Notation glossary (R, D, τ, κ, etc.)

Chapter 1 — Introduction

≈ 4,000 words

1.1 Motivation

1.1.1 Central banks as multi-instrument institutions
1.1.2 The fragmentation of the existing literature
1.1.3 Three canonical cases of policy dissonance
1.1.3.1 The 2013 Federal Reserve "taper tantrum"
1.1.3.2 The European Central Bank 2011 cycle
1.1.3.3 The 2023 US regional banking stress

1.2 Research Question

1.3 Approach in Brief

1.3.1 The structural decomposition
1.3.2 The dissonance metric
1.3.3 The fourteen-country empirical panel

1.4 Contributions

1.4.1 Theoretical
1.4.2 Empirical
1.4.3 Methodological

1.5 Scope and Limitations

1.5.1 The lens of monetary policy
1.5.2 What this thesis is not
1.5.3 The discretionary–systematic boundary

1.6 Structure of the Thesis

Chapter 2 — Literature Review

≈ 7,000 words

2.1 Frameworks for Central Bank Decomposition

2.1.1 The Svensson optimisation framework
2.1.2 Tinbergen and the multi-instrument tradition
2.1.3 The IMF Integrated Policy Framework
2.1.4 Borio and the BIS multi-mandate view

2.2 Reaction Functions

2.2.1 The Taylor rule and its descendants
2.2.2 Forward-looking and smoothed specifications
2.2.3 Open-economy and EM extensions
2.2.4 Estimated reaction functions for the panel countries

2.3 The Instrument Set

2.3.1 Forward guidance identification
2.3.2 Macroprudential policy measurement
2.3.3 Balance sheet policy and quantitative easing
2.3.4 FX intervention and capital flow management
2.3.5 Liquidity facilities

2.4 Transmission and Intermediary Frictions

- 2.4.1 The monetary transmission channels framework
- 2.4.2 Intermediary asset pricing
- 2.4.3 The dealer balance sheet and CIP basis
- 2.4.4 Sovereign-bank loops

2.5 Regime-Switching and Structural Break Detection

- 2.5.1 Markov-switching macroeconomic models
- 2.5.2 Structural break methods
- 2.5.3 Time-varying parameter VARs
- 2.5.4 Macro early-warning systems

2.6 Emerging Market Monetary Economics

- 2.6.1 Fear of floating and the trilemma
- 2.6.2 The global financial cycle
- 2.6.3 Sudden stops and capital flow reversals

2.7 Synthesis: Positioning the Contribution

- 2.7.1 What the literature decomposes
- 2.7.2 What the literature does not measure
- 2.7.3 The dissonance overlay

Chapter 3 — Theoretical Framework

≈ 6,000 words

3.1 The Structural Object

- 3.1.1 The directed graph G = (N, E)
- 3.1.2 The ten hierarchical levels
- 3.1.3 Tree, DAG, or general directed graph
- 3.1.4 Edge classes: transmission, feedback, peer
- 3.1.5 The pedagogical projection vs the formal object

3.2 Node Taxonomy

- 3.2.1 Agent nodes
- 3.2.2 Market nodes
- 3.2.3 Instrument nodes
- 3.2.4 Mathematical implications of species

3.3 Topology

- 3.3.1 Upstream and downstream sets
- 3.3.2 Transmission edges
- 3.3.3 Feedback edges
- 3.3.4 Peer edges and within-level relationships
- 3.3.5 The closed loop

3.4 The Consistency Model

- 3.4.1 The consistency function R_n
- 3.4.2 Dissonance as the violation of R_n
- 3.4.3 The normal-regime threshold τ_n
- 3.4.4 Theoretical grounding requirement

3.5 Dissonance Taxonomy

3.5.1 Intra-node dissonance
3.5.2 Inter-node dissonance
3.5.3 Vertical dissonance
3.5.4 Temporal dissonance

3.6 Country-Specific Parameterisation

3.6.1 Architecture universality
3.6.2 Sub-node activation
3.6.3 Numerical parameter specificity
3.6.4 Threshold non-linearities

3.7 Aggregation

3.7.1 The country-level composite $D_{c,t}$
3.7.2 Equal-weight scheme
3.7.3 Hierarchical scheme
3.7.4 PCA-extracted scheme
3.7.5 Theory-weighted scheme

3.8 Auxiliary Primitives

3.8.1 The modelling triad (state, dynamics, transmission)
3.8.2 The conditional transmission multiplier κ
3.8.3 Constraint binding
3.8.4 Frequency alignment
3.8.5 System regime characterisation

3.9 Summary of Definitions

Chapter 4 — Empirical Strategy

≈ 5,000 words

4.1 Sample Design

4.1.1 Country selection: the DM panel
4.1.2 Country selection: the EM panel
4.1.3 Time period and frequency
4.1.4 The DM/EM comparison as design feature

4.2 Node-Level Estimation

4.2.1 L1 reaction function estimation
4.2.2 L2 loss function and constraint indicators
4.2.3 L3 instrument dissonance metrics
4.2.4 L4 market dissonance metrics
4.2.5 L5 intermediary composites
4.2.6 L6 expectations gaps
4.2.7 L7 financial conditions composition

4.3 Composite Construction

4.3.1 Aggregation across nodes
4.3.2 Comparison of weighting schemes
4.3.3 Robustness across schemes as a result

4.4 Predictive Testing

4.4.1 Dependent variables: asset returns
4.4.2 Forecast horizons
4.4.3 Within-country regressions
4.4.4 Panel regressions
4.4.5 Out-of-sample design

4.5 Baseline Comparators

4.5.1 Statistical regime filters
4.5.2 Market-based stress indicators
4.5.3 Macro early-warning systems
4.5.4 Encompassing regression specification

4.6 Robustness Strategy

4.6.1 Drop-one-node analysis
4.6.2 Drop-one-country analysis
4.6.3 Alternative threshold calibrations
4.6.4 Crisis vs non-crisis sub-samples

4.7 Pre-Registration

Chapter 5 — Data

≈ 4,000 words

5.1 Database Architecture

5.1.1 Schema and storage
5.1.2 Country-time-frequency harmonisation
5.1.3 Version control

5.2 Data Sources by Node Level

5.2.1 Macroeconomic data (L1, L2, L10)
5.2.2 Central bank action data (L3)
5.2.3 Market data (L4, L9)
5.2.4 Intermediary data (L5)
5.2.5 Expectations data (L6)
5.2.6 Financial conditions (L7)

5.3 Data Quality Assessment

5.3.1 Coverage matrix by country and node
5.3.2 Frequency mismatches
5.3.3 Proxy decisions and justifications

5.4 Limitations and Scope Restrictions

5.4.1 NBFI data in emerging markets
5.4.2 Dealer-balance-sheet measurement
5.4.3 Communication data in non-English-language regimes

Chapter 6 — South Africa Pilot

≈ 5,000 words

6.1 Rationale for South Africa as Pilot

6.2 Node-by-Node Implementation

6.2.1 Reaction function: Aron-Muellbauer specification
6.2.2 Instrument vector: SARB historical record
6.2.3 Communication: NLP on MPC statements
6.2.4 Intermediaries: the Big Four banks and the sovereign
6.2.5 Markets: the R-bond curve and ZAR

6.3 The South African Dissonance Time Series, 2005–2024

6.3.1 Descriptive overview
6.3.2 Identified regime events
6.3.3 Composite signal under each weighting scheme

6.4 Sanity Checks

6.4.1 The 2008–09 global financial crisis
6.4.2 The 2013 taper tantrum
6.4.3 The 2015–16 commodity cycle
6.4.4 The 2020 COVID shock
6.4.5 The 2022–23 hiking cycle

6.5 Preliminary Predictive Analysis

6.6 Framework Adjustments Pre-Panel

6.6.1 Implementation choices not in the pre-registration
6.6.2 Justifications and audit trail

Chapter 7 — Panel Results

≈ 8,000 words

7.1 Cross-Country Dissonance Patterns

7.1.1 Descriptive statistics by country
7.1.2 Node-level firing frequencies
7.1.3 Common factor analysis across the panel

7.2 The Developed Market Sub-Panel

7.3 The Emerging Market Sub-Panel

7.4 Direct DM–EM Comparison

7.4.1 Frequency and magnitude differences
7.4.2 Node-level firing differences
7.4.3 Asset-class response differences

7.5 Predictive Regressions

7.5.1 Local-currency sovereign bond returns
7.5.2 Exchange rate returns
7.5.3 Equity index returns
7.5.4 Sovereign CDS returns

7.6 Out-of-Sample Performance

7.6.1 Expanding-window scheme

7.6.2 Comparison to in-sample fit

7.7 Encompassing Regressions Against Baselines

7.7.1 Against Markov-switching regime probability

7.7.2 Against FCI z-scores

7.7.3 Against the excess bond premium analogue

7.7.4 Against the Kaminsky-Reinhart signals approach

7.8 Robustness

7.8.1 Across weighting schemes

7.8.2 Drop-one-node

7.8.3 Drop-one-country

7.8.4 Alternative thresholds

Chapter 8 — Mechanism and Case Studies

≈ 5,000 words

8.1 Conditional Sharpe Analysis

8.1.1 Equilibrium-regime performance

8.1.2 Transition-regime performance

8.1.3 The regime-change-detector hypothesis

8.2 Node-Level Attribution

8.2.1 Which nodes carry the most predictive weight

8.2.2 Decomposition of $D_{c,t}$ contribution to forecast accuracy

8.3 Asset-Level Attribution

8.3.1 Mapping fired nodes to asset responses

8.3.2 Differentiating across asset classes

8.4 Case Studies

8.4.1 The 2013 US taper tantrum

8.4.2 The ECB 2011 cycle

8.4.3 The 2023 US regional banking stress

8.4.4 Turkey 2018

8.4.5 South Africa 2020 and 2022–23

8.4.6 A null case: an event where dissonance did not predict

8.5 Limits of the Framework

8.5.1 What the framework misses

8.5.2 False positives and false negatives

Chapter 9 — Discussion and Conclusion

≈ 3,000 words

9.1 Summary of Findings

9.2 Contributions Revisited

9.2.1 Theoretical contribution

9.2.2 Empirical contribution
9.2.3 Methodological contribution

9.3 Implications for Monetary Economics

9.4 Implications for Empirical Asset Pricing

9.5 Implications for Investment Practice

9.6 Limitations

9.7 Future Research

9.7.1 Extending to developed market intermediary detail
9.7.2 Structural interpretation as monetary shock identification
9.7.3 Cross-level dissonance as a standalone research programme
9.7.4 Operationalisation for live portfolio management

Back Matter

References

Appendix A — Mathematical Definitions and Proofs

Appendix B — Data Sources, Transformations, and Coverage Matrix

Appendix C — Pre-Registration Document (timestamped)

Appendix D — Country-Specific Reaction Function Estimates

Appendix E — Node-Level Dissonance Time Series, Full Panel

Appendix F — Additional Robustness Tables

Appendix G — Code Repository and Replication Instructions

Structural choices worth flagging

1. Chapter 6 (SA Pilot) stands alone as a chapter. This treats the pilot as a substantive deliverable in its own right rather than burying it in the appendix. The advantage is a complete case study that can be excerpted as a standalone paper. The cost is some redundancy with Chapter 7.

2. Chapter 8 includes a deliberate null case study. Asymmetric reporting — only showing cases where the framework "works" — would weaken the thesis. Including a case where dissonance did not predict (or fired falsely) is what makes the empirical claim credible.

3. Appendix C is the pre-registration document, timestamped. Unusual for an economics thesis but it's the cleanest way to defend against post-hoc-tinkering critiques and pre-empt the most likely examiner concern.

> *Next: Section 3.1 — the Structural Object — where the directed-graph-vs-tree question is formally resolved.*